

Lecture 6 – ANOVA and Linear Models

STAT/BIOF/GSAT 540

Keegan Korthauer

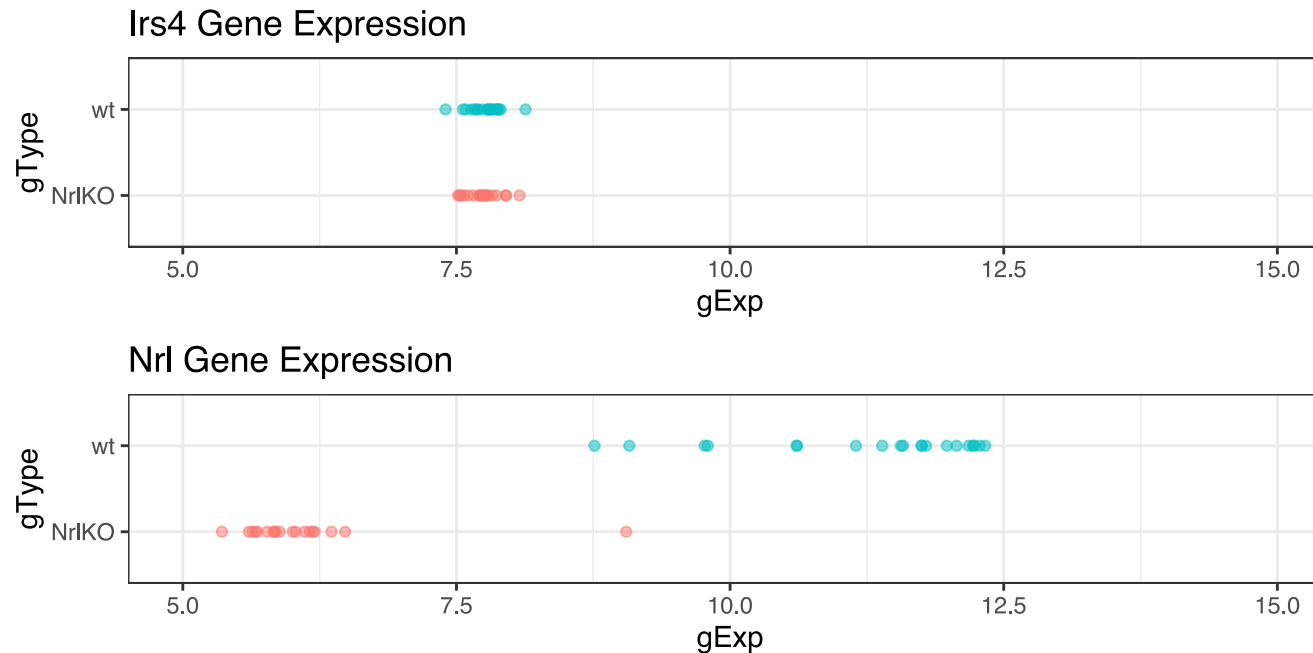
2020/01/22

Slides by: Gabriela Cohen Freue with contributions from Jenny Bryan and Keegan Korthauer


Are these genes truly different in NrlKO compared to WT?

H_0 : the expression level of gene A is the same in both conditions.

Is there **enough** evidence in the data to reject H_0 ?



Statistics: use a random sample to learn about the population



$$Y \sim F$$

$$Z \sim G$$

$$E[Y] = \mu_Y$$

$$E[Z] = \mu_Z$$

$$Y_1, Y_2, \dots, Y_{n_Y}$$

$$Z_1, Z_2, \dots, Z_{n_Z}$$

$$\hat{\mu}_Y = \bar{Y} = \frac{\sum_{i=1}^{n_Y} Y_i}{n_Y}$$

$$\hat{\mu}_Z = \bar{Z} = \frac{\sum_{i=1}^{n_Z} Z_i}{n_Z}$$

$$H_0: \mu_Y = \mu_Z$$

$$\bar{Y} = \bar{Z}$$

Last class: hypothesis testing

1. Define a **test statistic** to test

- 2-sample t -test
- Welch t -test
- Wilcoxon rank-sum test
- Kolmogorov-Smirnov test

2. Compute the **observed value** for the test statistic

3. Compute the probability of seeing a test statistic as extreme as that observed, under the **null sampling distribution** (p-value)

4. Make a decision about the **significance** of the results, based on a pre-specified value (alpha, significance level)

We can run these tests in R

Example: use the `t.test` function to test using a classical 2-sample *t*-test.

```
miniDat %>% subset(gene == "Irs4") %>% t.test(gExp ~ gType, data = .,  
  var.equal = TRUE)
```

```
##  
##      Two Sample t-test  
##  
## data:  gExp by gType  
## t = -0.52865, df = 37, p-value = 0.6002  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
##  -0.12597002  0.07383844  
## sample estimates:  
## mean in group NrlK0      mean in group wt  
##           7.739684           7.765750
```

Today...

- show how to compare means of different groups (2 or more) using a linear regression model
 - dummy variables to model the levels of a qualitative explanatory variable
- write a linear model using matrix notation
 - understand which matrix is built by R
- distinguish between conditional and marginal effects
 - β -tests vs γ -tests

```
> t.test(gExp ~ gType, miniDat,  
+       subset = gene == "Irs4", var.equal = TRUE)
```

two sample t test

$$H_0 : \mu_1 = \mu_2$$

```
> summary(aov(gExp ~ gType, miniDat,  
+            subset = gene == "Irs4"))
```

(one-way) analysis of variance
“ANOVA”

```
> summary(lm(gExp ~ gType, miniDat,  
+           subset = gene == "Irs4"))
```

linear model
linear regression

It seems that we can use any of these methods to test


```
> t.test(gExp ~ gType, miniDat,  
+       subset = gene == "Irs4", var.equal = TRUE)
```

Two Sample t-test

```
data: gExp by gType  
t = 0.5286, df = 37, p-value = 0.6002  
<snip, snip>  
sample estimates:  
mean in group wt mean in group NrlKO  
7.765750 7.739684
```

```
> summary(aov(gExp ~ gType, miniDat,  
+       subset = gene == "Irs4"))  
          Df Sum Sq Mean Sq F value Pr(>F)  
gType      1  0.0066  0.00662    0.279    0.6  
Residuals 37  0.8764  0.02369
```

```
> summary(lm(gExp ~ gType, miniDat,  
+       subset = gene == "Irs4"))  
<snip, snip>  
Coefficients:
```



$$7.739684 - 7.765750 = -0.026066$$

$$-0.5286494^2 = 0.2794702$$

t-test vs linear regression: why the same results?


```
irs4Dat <- subset(miniDat, gene=="Irs4")  
ttest.irs4 <- t.test(gExp ~ gType, irs4Dat, var.equal = TRUE)  
list("t value"=ttest.irs4$stat, "p-value"=ttest.irs4$p.value)
```

```
## $t value  
##      t  
## -0.5286494  
##  
## $p-value  
## [1] 0.6002058
```

```
lm.irs4 <- summary(lm(gExp ~ gType, irs4Dat))  
list("t value"=lm.irs4$coeff[2,3], "p-value"=lm.irs4$coeff[2,4])
```

```
## $t value  
## [1] 0.5286494  
##  
## $p-value  
## [1] 0.6002058
```

t -test *vs* linear regression: where's the line?



Note that the y -axis in these plots is not numerical, thus a line in this space does not have any mathematical meaning.

Why can we run a t -test with a **linear** regression model?

From t -test to linear regression

Let's change the notation to give a common framework to all methods



We can use a subindex to distinguish observations from each group, i.e.,

where y_{ij} is the response for the j th observation in the i th group, or y_i identifies the groups; and j identifies the observations within each group

For example: y_{11} is the first observation in group 1 or WT

The goal is to test

using data from the model

where μ_j or σ_j ; and μ_j .

For simplicity, we assume a common distribution for all groups

Note that the population means are given by μ_j , i.e., the model is written with a cell-means - parametrization

Note that for each group, the **population** mean is given by

A natural **estimator** of the population mean is the **sample mean**

Classical hypothesis testing methods use the group sample means as estimators

See, for example, the `t.test` function in R:

```
ttest.irs4$estimate
```

```
## mean in group Nr1K0    mean in group wt  
##           7.739684           7.765750
```

However, the `lm` function reports other estimates, **why?**

```
(means.irs4 <- as.data.frame(irs4Dat %>% group_by(gType) %>%  
  summarize(meanGroups = mean(gExp, digits = 6))))
```

```
##   gType meanGroups  
## 1 NrlKO    7.739684  
## 2   wt     7.765750
```

```
lm.irs4$coefficients[,1]
```

```
## (Intercept)      gTypewt  
##  7.73968421  0.02606579
```

↓

(Intercept) is the **sample mean** of NrlKO group

but gTypewt is **not** the sample mean of the WT group

Parametrizations: which parameters we use to write the model?

By default, the `lm` does not use the cell-means parametrization. The goal is to *compare* the means, not to study each in isolation.

From **cell-means** - :




to **reference-treatment effect** - :

Note that for each group, the population mean is given by

and

compares the means

Relation between parametrizations



lm reports the sample mean of the **reference** group (Nr1KO):

and the **treatment effect**, i.e., difference between the sample means of both groups:


```
lm.irs4$coefficients[, 1]
```

```
## (Intercept)      gTypewt  
## 7.73968421 0.02606579
```

```
data.frame(meanWT = means.irs4[1, 2],  
           meanDiff = diff(means.irs4$meanGroups))
```

```
##      meanWT  meanDiff  
## 1 7.739684 0.02606579
```

We still haven't answered ... where's the line??



Dummy variables

Let's re-write our model using **dummy** (or indicator) variables:



Note that β_0 , because β_1 and β_2 , because
(for all i)

The second form is written as a **linear** (multiple) regression, with a special (**dummy**) explanatory variable

Using a dummy variables to model our categorical variables `gt type` we can perform a **2-sample *t*-test** with a linear model

```
list("t value"=ttest.irs4$stat,"p-value"=ttest.irs4$p.value)
```

```
## $t value  
##          t  
## -0.5286494  
##  
## $p-value  
## [1] 0.6002058
```

```
list("t value"=lm.irs4$coeff[2,3],"p-value"=lm.irs4$coeff[2,4])
```

```
## $t value  
## [1] 0.5286494  
##  
## $p-value  
## [1] 0.6002058
```


Beyond 2-groups comparisons: difference of means

“cell-means”


$$Y_{ij} = \mu_j + \varepsilon_{ij}$$

“reference-treatments”

$$Y_{ij} = \theta + \tau_j + \varepsilon_{ij}, (\tau_1 = 0)$$



More than 2
groups!



Dummy variables can be used to model one *or more* categorical variables with 2 *or more* levels!

2-sample *t*-test using a linear model

1-way ANOVA with many levels using a linear model

This is why R can estimate all of them with `lm()`

in general, yet another parametrization is used to present ANOVA

t-test

Special case of **ANOVA**, but with ANOVA you can compare **more than two groups** and **more than one factor**.

ANOVA

Special case of **linear regression**, but with linear regression you can include **quantitative variables** in the model.

Linear regression

Provides a unifying framework to model the association between a response **many quantitative and qualitative variables**.

In R: all can be computed using the `lm ()` function.

Linear models using matrix notation

the column vector of the responses
one element per experimental unit

a column vector
of the errors



The diagram shows the equation $Y = X\alpha + \epsilon$ with four arrows pointing to its components: one from the text 'the column vector of the responses' to Y , one from 'a column vector of the errors' to ϵ , one from 'a (design) matrix that represents covariate info' to X , and one from 'a column vector of the parameters in the linear model' to α .

$$Y = X\alpha + \epsilon$$

a (design) matrix that represents covariate
info, one row per experimental unit

a column vector of the parameters in the
linear model

It will become handy to write our model using matrix notation

Let's form an matrix for a 3-groups comparison:

$$\begin{bmatrix} Y_{11} \\ \vdots \\ Y_{n_1 1} \\ Y_{12} \\ \vdots \\ Y_{n_2 2} \\ Y_{13} \\ \vdots \\ Y_{n_3 3} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ \vdots & \vdots & \vdots \\ 1 & 0 & 0 \\ 1 & 1 & 0 \\ \vdots & \vdots & \vdots \\ 1 & 1 & 0 \\ 1 & 0 & 1 \\ \vdots & \vdots & \vdots \\ 1 & 0 & 1 \end{bmatrix} \begin{bmatrix} \theta \\ \tau_2 \\ \tau_3 \end{bmatrix} + \begin{bmatrix} \varepsilon_{11} \\ \vdots \\ \varepsilon_{n_1 1} \\ \varepsilon_{12} \\ \vdots \\ \varepsilon_{n_2 2} \\ \varepsilon_{13} \\ \vdots \\ \varepsilon_{n_3 3} \end{bmatrix}$$

response Y
design matrix X
regression parameters
error term

$$Y = X\alpha + \varepsilon$$

Note that θ and τ_1 become the 1st and 2nd columns of X : θ for the reference group; τ_1 for the 2nd group; and τ_2 for the 3rd group

$$Y_{ij} = \theta + \tau_2 \times x_{ij2} + \tau_3 \times x_{ij3} + \varepsilon_{ij}$$

$$\begin{bmatrix} \underline{Y_{11}} \\ \vdots \\ Y_{n_1 1} \\ \underline{Y_{12}} \\ \vdots \\ Y_{n_2 2} \\ \underline{Y_{13}} \\ \vdots \\ Y_{n_3 3} \end{bmatrix} = \begin{bmatrix} \underline{1} & 0 & 0 \\ \vdots & \vdots & \vdots \\ 1 & 0 & 0 \\ \underline{1} & 1 & 0 \\ \vdots & \vdots & \vdots \\ 1 & 1 & 0 \\ \underline{1} & 0 & 1 \\ \vdots & \vdots & \vdots \\ 1 & 0 & 1 \end{bmatrix} \begin{bmatrix} \theta \\ \tau_2 \\ \tau_3 \end{bmatrix} + \begin{bmatrix} \underline{\varepsilon_{11}} \\ \vdots \\ \varepsilon_{n_1 1} \\ \underline{\varepsilon_{12}} \\ \vdots \\ \varepsilon_{n_2 2} \\ \underline{\varepsilon_{13}} \\ \vdots \\ \varepsilon_{n_3 3} \end{bmatrix}$$

Note that

Note that

Note that

Which is the same as

Which is the same as

$$Y = X\alpha + \varepsilon$$

$$\begin{bmatrix} Y_{11} \\ Y_{21} \\ \vdots \\ Y_{n33} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ \vdots & \vdots & \vdots \\ 1 & 0 & 0 \\ 1 & 1 & 0 \\ \vdots & \vdots & \vdots \\ 1 & 1 & 0 \\ 1 & 0 & 1 \\ \vdots & \vdots & \vdots \\ 1 & 0 & 1 \end{bmatrix} \begin{bmatrix} \theta \\ \tau_2 \\ \tau_3 \end{bmatrix} + \begin{bmatrix} \varepsilon_{11} \\ \varepsilon_{21} \\ \vdots \\ \varepsilon_{n33} \end{bmatrix}$$

Diagram illustrating the reference-treatment parametrization:

- Reference group: μ_1 (indicated by a red box and arrow pointing to θ)
- Treatment 2 effect: $\mu_2 - \mu_1$ (indicated by a blue box and arrow pointing to τ_2)
- Treatment 3 effect: $\mu_3 - \mu_1$ (indicated by a green box and arrow pointing to τ_3)

Note that the model is still written with a reference-treatment parametrization (difference of means)

Linear regression can include quantitative & qualitative covariates.

$$Y = X\alpha + \varepsilon$$

This gives us a VERY FLEXIBLE framework!!

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 1 & 0 \\ 1 & 0 & 0 & 1 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 0 & 1 \end{bmatrix}$$

1 categorical
covariate

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 1 & 1 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 1 & 1 & 1 \end{bmatrix}$$

2 categorical
covariates

$$\begin{bmatrix} 1 & 1.22 \\ 1 & 2.02 \\ 1 & 1.42 \\ \vdots & \vdots \\ 1 & 1.89 \\ 1 & 2.01 \\ \vdots & \vdots \\ 1 & 1.56 \\ 1 & 2.17 \\ 1 & 1.51 \end{bmatrix}$$

1 continuous
covariate

$$\begin{bmatrix} 1 & 0 & 1.22 & 0 \\ 1 & 0 & 2.02 & 0 \\ 1 & 0 & 1.42 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 1.89 & 0 \\ 1 & 1 & 2.01 & 2.01 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 1 & 1.56 & 1.56 \\ 1 & 1 & 2.17 & 2.17 \\ 1 & 1 & 1.51 & 1.51 \end{bmatrix}$$

1 continuous
1 categorical

AND MANY MORE

Tip: ?model.matrix

Linear in the parameters : can contain , , etc.

How it works in practice using `lm()` in R



```
lm(y ~ x, data = yourData)
```

`y ~ x`: formula,
 `y` numeric,
`x` numeric and/or factor

`yourData`: data.frame in which x and y are
to be found (optional but recommended)

By default, R uses a ref-tx parametrization but you can control that!

- Mathematically, X is a numeric matrix
- If your data contains categorical variables (e.g., `gType`), you need to set them as **factors**
- R creates appropriate dummy variables for factors!

```
str(irs4Dat$gType)
```

```
## Factor w/ 2 levels "NrlK0","wt": 2 2 2 2 1 1 1 2 2 2 ...
```

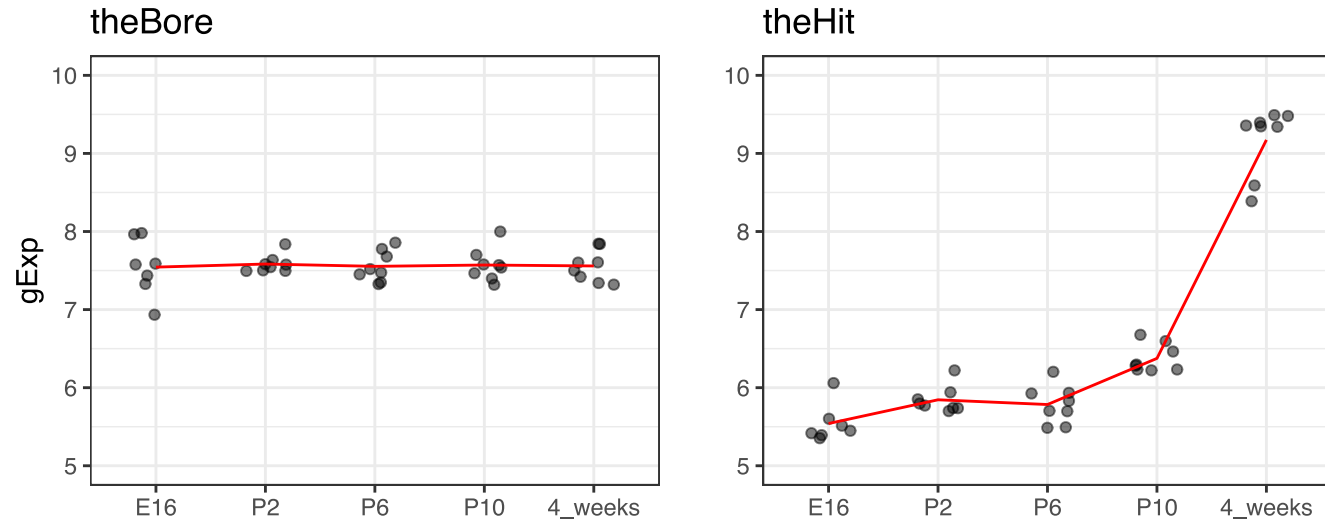
Under the hood, R creates a numeric :

```
data.frame(X = model.matrix(gExp ~ gType, irs4Dat),  
           gType = irs4Dat$gType) %>% head(10)
```

```
##      X..Intercept. X.gTypewt gType  
## 1             1         1    wt  
## 2             1         1    wt  
## 3             1         1    wt  
## 4             1         1    wt  
## 5             1         0 Nr1KO  
## 6             1         0 Nr1KO  
## 7             1         0 Nr1KO  
## 8             1         1    wt  
## 9             1         1    wt  
## 10            1         1    wt
```

Beyond 2-group comparisons in our case study:

Is the expression of gene A the same at all developmental stages?

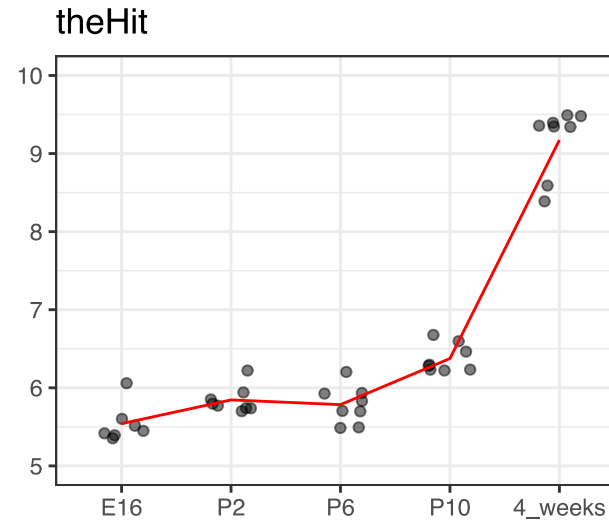



Note: 4W = 4_weeks

The **sample** means:


```
with(devDat, tapply(gExp, list(devStage, gene), mean))
```

```
##           theBore  theHit  
## E16      7.544143 5.540857  
## P2       7.583500 5.844875  
## P6       7.554000 5.784250  
## P10      7.571000 6.375125  
## 4_weeks  7.559000 9.173375
```



"theHit" with significant time ("treatment") effect

	cellMeans	txEffects
E16	5.540857	0.0000000
P2	5.844875	0.3040179
P6	5.784250	0.2433929
P10	6.375125	0.8342679
4_weeks	9.173375	3.6325179



"theHit" with significant time ("treatment") effect

Can you guess the size of the matrix??

- | How many dummy variables do we need?

"theHit" with significant time ("treatment") effect

```
##   devStage cellMeans txEffects
## 1      E16  5.540857 0.0000000
## 2      P2  5.844875 0.3040179
## 3      P6  5.784250 0.2433929
## 4     P10  6.375125 0.8342679
## 5 4_weeks  9.173375 3.6325179
```

We need 4 dummy variables to estimate and test 4 time differences:

| : P2 vs E16, : P6 vs E16, : P10 vs E16, : 4W vs E16)

Mathematically:

Notation: , where is an index for the observation, for the level of devStage, and for the name of the dummy variable

Under the hood, R creates a numeric :

```
X.matrix <- data.frame(X = model.matrix(gExp ~ devStage, irs4Dat),  
  devStage = irs4Dat$devStage)
```

##	X..Intercept.	X.dStP2	X.dStP6	X.dStP10	X.dS4W	dS
## 1	1	0	0	0	0	E16
## 2	1	0	0	0	0	E16
## 3	1	0	0	0	0	E16
## 4	1	0	0	0	0	E16
## 5	1	0	0	0	0	E16
## 6	1	0	0	0	0	E16
## 7	1	0	0	0	0	E16
## 8	1	1	0	0	0	P2
## 9	1	1	0	0	0	P2
## 10	1	1	0	0	0	P2
## 11	1	1	0	0	0	P2
## 12	1	1	0	0	0	P2
## 13	1	1	0	0	0	P2
## 14	1	1	0	0	0	P2
## 15	1	1	0	0	0	P2
## 16	1	0	1	0	0	P6

Note: column names changed and first 16 rows displayed to fit output in the page (code hidden)

```
summary(lm(gExp~devStage,subset(devDat,gene=="theHit")))$coeff
```

##		Estimate	Std. Error	t value	Pr(> t)
##	(Intercept)	5.5408571	0.1021381	54.248698	1.307554e-34
##	devStageP2	0.3040179	0.1398583	2.173756	3.678022e-02
##	devStageP6	0.2433929	0.1398583	1.740282	9.085489e-02
##	devStageP10	0.8342679	0.1398583	5.965093	9.559065e-07
##	devStage4_weeks	3.6325179	0.1398583	25.972843	5.266481e-24

```
means.dev %>% mutate(txEffects=cellMeans-cellMeans[1])
```

##	devStage	cellMeans	txEffects
## 1	E16	5.540857	0.0000000
## 2	P2	5.844875	0.3040179
## 3	P6	5.784250	0.2433929
## 4	P10	6.375125	0.8342679
## 5	4_weeks	9.173375	3.6325179

Estimate:

or

we are not usually interested in testing this hypothesis

```
summary(lm(gExp~devStage,subset(devDat,gene=="theHit")))$coeff
```

##		Estimate	Std. Error	t value	Pr(> t)
##	(Intercept)	5.5408571	0.1021381	54.248698	1.307554e-34
##	devStageP2	0.3040179	0.1398583	2.173756	3.678022e-02
##	devStageP6	0.2433929	0.1398583	1.740282	9.085489e-02
##	devStageP10	0.8342679	0.1398583	5.965093	9.559065e-07
##	devStage4_weeks	3.6325179	0.1398583	25.972843	5.266481e-24

```
means.dev %>% mutate(txEffects=cellMeans-cellMeans[1])
```

##	devStage	cellMeans	txEffects
## 1	E16	5.540857	0.0000000
## 2	P2	5.844875	0.3040179
## 3	P6	5.784250	0.2433929
## 4	P10	6.375125	0.8342679
## 5	4_weeks	9.173375	3.6325179

Estimate:

or

we *are* usually interested in testing this hypothesis: first 2 days after birth

```
summary(lm(gExp~devStage,subset(devDat,gene=="theHit")))$coeff
```

##		Estimate	Std. Error	t value	Pr(> t)
##	(Intercept)	5.5408571	0.1021381	54.248698	1.307554e-34
##	devStageP2	0.3040179	0.1398583	2.173756	3.678022e-02
##	devStageP6	0.2433929	0.1398583	1.740282	9.085489e-02
##	devStageP10	0.8342679	0.1398583	5.965093	9.559065e-07
##	devStage4_weeks	3.6325179	0.1398583	25.972843	5.266481e-24


```
means.dev %>% mutate(txEffects=cellMeans-cellMeans[1])
```

##	devStage	cellMeans	txEffects
## 1	E16	5.540857	0.0000000
## 2	P2	5.844875	0.3040179
## 3	P6	5.784250	0.2433929
## 4	P10	6.375125	0.8342679
## 5	4_weeks	9.173375	3.6325179

Estimate:

or

we *are* usually interested in
testing this hypothesis: 4 weeks
after birth



We generally test two types of null hypotheses:

VS

for each j **individually**

e.g., Is gene A differentially expressed 2 days after birth?

VS

for all j **at the same time**

e.g., Is gene A significantly affected by time (devStage)?

Two types of null hypotheses in R:

$$Y = X\alpha + \varepsilon$$
$$\alpha = (\theta, \tau_{P2}, \tau_{P6}, \tau_{P10}, \tau_{4_weeks})$$

$$H_0 : \tau_j = 0$$

vs

$$H_0 : \tau_j \neq 0$$

for each j individually

$$H_0 : \tau_j = 0 \quad \text{AND statement}$$

vs

$$H_0 : \tau_j \neq 0 \quad \text{OR statement}$$

for all j at the same time

```
> summary(hitFit)
Call:
lm(formula = gExp ~ devStage, <blah, blah>)
<snip, snip>
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    5.5409      0.1021  54.249  < 2e-16 ***
devStageP2      0.3040      0.1399   2.174   0.0368 *
devStageP6      0.2434      0.1399   1.740   0.0909 .
devStageP10     0.8343      0.1399   5.965  9.56e-07 ***
devStage4_weeks 3.6325      0.1399  25.973  < 2e-16 ***
---
<snip, snip>
F-statistic: 243.4 on 4 and 34 DF, p-value: < 2.2e-16
```

F -test and overall significance of one or more covariates

- the t -test in linear regression allows us to test single hypotheses:
- but we often like to test multiple hypotheses *simultaneously*:

the F -test allows us to test such compound tests

To conclude

- we can use different parametrizations to write statistical models

From **cell-means** - :

to **reference-treatment effect** - : (used by default by `lm`)

- we can compare group means (2 or more) using a linear model
 - **dummy variables** (e.g.,) to model the levels of a qualitative explanatory variables
 - qualitative variables need to be set as "factors" in the data --> R creates the dummy variables

- we can write a linear model using matrix notation:

- **Linear model** can include **quantitative & qualitative covariates**.

$$Y = X\alpha + \varepsilon$$

This gives us a VERY FLEXIBLE framework!!

$\begin{bmatrix} 1 & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 0 & 1 \end{bmatrix}$ <p>1 categorical covariate</p>	$\begin{bmatrix} 1 & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 1 & 1 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 1 & 1 & 1 \end{bmatrix}$ <p>2 categorical covariates</p>	$\begin{bmatrix} 1 & 1.22 \\ 1 & 2.02 \\ 1 & 1.42 \\ \vdots & \vdots \\ 1 & 1.89 \\ 1 & 2.01 \\ \vdots & \vdots \\ 1 & 1.56 \\ 1 & 2.17 \\ 1 & 1.51 \end{bmatrix}$ <p>1 continuous covariate</p>	$\begin{bmatrix} 1 & 0 & 1.22 & 0 \\ 1 & 0 & 2.02 & 0 \\ 1 & 0 & 1.42 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 0 & 1.89 & 0 \\ 1 & 1 & 2.01 & 2.01 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & 1 & 1.56 & 1.56 \\ 1 & 1 & 2.17 & 2.17 \\ 1 & 1 & 1.51 & 1.51 \end{bmatrix}$ <p>1 continuous 1 categorical</p>
---	--	--	--

AND MANY MORE

Tip: ?model.matrix

- distinguish between single and joint hypotheses:

- -tests vs -tests